

Bitcode

Measured Play — Reputation, Progression & Mandates

The positive game layer: making honest contribution compelling · every mechanic grounded in real measurements · DRAFT for founder review.

1 · Why a game layer, and why ours is different

Open platforms live and die on the 90-9-1 rule: most watch, few contribute, almost no one carries the weight. The classic fix is gamification; the classic failure is Goodhart — reward a number and people farm the number. Bitcode is uniquely positioned to avoid that failure, because it already runs the most expensive part of any game: an honest scoring engine. The measurement stack produces statuses, volumes, fits, settlements, and proof roots as a matter of law. The game layer invents no new scores — it renders the ones that already exist, and attaches progression, identity, and delight to them.

The prime rule: no mechanic may mint, price, or measure. Gamification reads the measurement substrate; it never writes it. Anything that would reward raw volume is presumed an attack vector (see Anti-Gaming, layer D) until proven otherwise.

2 · Seller progression — the measured ladder

Levels are earned states computed from measurement facts, not points. Illustrative ladder (names and thresholds are product decisions):

Level	Earned by (measurement facts)	Unlocks
Depositor	First admitted pack, catalog-complete	Profile, measured badge
Measured Seller	N packs at deep/thin measure mode; zero gate trips	Depository feature placement
Proven Seller	First mandate won (settled read); realized need-fit above threshold	Earnings dashboard, priority measurement
Founding Supplier	Sustained mandates won; breadth across kinds or domains	Founding-supply terms, roadmap voice

Badges are measurement facts wearing a name. “Deep-Measured ×10” (ten packs at deep mode), “Cleanroom” (zero hygiene-gate trips across all deposits), “First Settle,” “Original” (originality above threshold when that absolute lands), “Conviction” (cumulative \$BTD elected over cash). Every badge is auditable — click it, see the readings that earned it. A badge no one can fake is the only kind worth displaying.

3 · Boards and indices — leaderboards that inform the market

Bitcode's leaderboards double as market data. Ranked, opt-in surfaces: most-measured sellers, deepest coverage, top realized need-fit, domain boards (best-measured packs per knowledge domain), and a conviction board (share of payout elected as \$BTD — the revealed-preference signal from tokenomics made social). Company dashboards serve the IP-recycling pool: an enterprise sees its own idle-code estate becoming measured, listed, and earning — progress bars over repositories, not vanity counters. The same rendering that motivates a solo researcher is a monetization report for a CTO.

4 · Mandates — demand-shaped contribution

A mandate is real market language — an instruction from capital to go get something done — and that is exactly what an unmet Need is. Sellers don't complete tasks on Bitcode; they win mandates.

Open mandates. Unmet or under-served Needs become open mandates on a public board: “this demand exists, no pack fits it well yet.” The seller whose pack best fits wins the mandate — earning the settlement plus the mandate mark on their record. This routes contributor energy toward real demand, the opposite of volume farming.

Depth mandates. Standing mandates from the exchange itself: raise a pack's measure mode from thin to deep; add the sibling tests that unlock verification kinds when they land. Fulfilling one improves measurement quality — the instrument pays for its own sharpening.

Estate mandates (enterprise). For IP-recycling partners: staged mandates over their repository estate — scope, rubric-sign, deposit, first settle — the design-partner pilot rendered as a mandate line with progress the whole team can see.

Streaks, gently. Consecutive measured deposits or maintained packs (re-measured after reference bumps) keep supply fresh; streaks decay softly rather than punishing absence. No daily-login theater.

5 • Buyer-side reputation (light)

Buyers accrue a quiet counterpart: well-specified Needs (schema-complete, measurable), settled reads, and dispute-free history. It is surfaced to the market as reliability, not rank — sellers deciding which mandates to pursue deserve the signal. No buyer leaderboards; demand should feel like a counterparty, not a contest.

6 • Guardrails — the anti-Goodhart contract

Rule	Rationale
Read-only over the substrate	Mechanics render measurements; they never feed back into volumes, weights, fits, or mint (the fenced mint stays fenced)
No raw-volume rewards	Pack-count alone earns nothing; every mechanic requires a quality co-signal (measure mode, gate cleanliness, settlement, fit)
Auditable badges only	Every earned state links to the readings behind it — the honesty-status system extended to reputation
Opt-in visibility	Earnings and conviction boards are opt-in; enterprises default private
Red-team the mechanics	Each new mechanic is reviewed against the Anti-Gaming attack catalog before launch (does it subsidize D1–D7?)
Sybil-aware	Ladder states derive from material-level facts (fingerprint-clustered), not identity counts

7 • Potential Rollout

V1 costs almost nothing: badges and the measured ladder are queries over data the pipeline already writes (statuses, modes, gate results, settlements). V2 adds boards and company dashboards — rendering plus opt-in policy. V3 opens the mandate board, which wants the open Needs surface, and depth mandates, which want the verification kinds. Sequence follows the measurement roadmap; the game layer is always one step behind the instrument, never ahead of it.