

Bitcode

The IP-Recycling Thesis — and Supply-Side Strategy

Why a successful company can profitably sell its non-moat code — the largest, most defensible, hardest-to-unlock supply on Earth.

1. The supply landscape in one view

Three seller pools feed Bitcode: **(A) IP-recycling** — successful companies monetizing non-moat code; **(B) research-for-cash** — experts and hackers producing code for pay; **(C) defunct-IP** — dead or pivoted companies' abandoned code. One engine unlocks all three: **automated, source-agnostic measurement**. Bitcode accepts any code from anyone and prices it by machine, so supply scales like GitHub — not like a staffing agency.

Ranked by potential (not by ease): **IP-recycling is the heaviest** — the biggest, highest-quality, and only genuinely exclusive pool — but the hardest to prove and unlock. Research-for-cash is the fastest to ignite (self-serve, proven demand). Defunct-IP is the most uncontested. This memo centers the hard one, because proving it unlocks the whole prize.

2. The IP-Recycling Thesis

The claim. For a successful business, selling a non-trivial fraction of its *non-moat* code is positive-ROI. Keeping every line private leaves money on the table for no strategic gain. This is not about the crown jewels — the differentiating algorithms you should never sell — but about the large, ordinary remainder of a real codebase.

2.1 Anatomy of a codebase — most code is not the moat

Every codebase is layered:

- **Crown jewels** (a small fraction) — the differentiating logic that is the moat. Never sold.
- **Non-moat functional** (the majority) — infrastructure, integrations, tooling, data pipelines, auth and billing, utilities, glue code, test harnesses, framework wrappers.
- **Dead / deprecated** (a meaningful slice) — pivoted-away features, old versions, killed experiments.

By any engineer's experience, the crown jewels are a small share; the non-moat and dead layers are the bulk. Your competitors already have equivalent plumbing — it is not why you win, and exposing it (obfuscated) does not erode your advantage.

2.2 The arbitrage — a liability inside, an asset outside

Non-moat and dead code has **near-zero private value**: it sits idle and costs money to maintain. Industry technical debt runs on the order of **\$2.4T a year**, and firms spend up to **40% of IT budgets** just keeping legacy alive. Inside the company it is a liability. Outside, the same code is a high-value asset: real, production, rights-clean code is the scarcest and most-wanted AI training data — code lifts model reasoning, and the public web runs dry around 2028. The identical code is a cost inside and revenue outside. Bitcode captures that spread.

2.3 The ROI, by construction

Return = (market value: cash + \$BTD + ongoing co-ownership earnings) – (marginal cost ≈ 0 , the code is already written) – (exposure risk ≈ 0 with obfuscation and non-moat selection) – (competitive erosion ≈ 0 for non-moat code). Every subtraction rounds to zero; the market value is positive. ROI is positive by construction. Even if only **10–20%** of a large codebase is safely sellable, that is a recurring, near-zero-cost revenue line from an asset that today earns \$0 and costs maintenance.

2.4 Objection-handling — the defensibility

- **"It helps competitors."** Non-moat code confers no advantage — rivals already have equivalents. The buyers are AI labs training models, not competitors cloning your product. Obfuscation hides anything sensitive, and you choose what to deposit.
- **"It's a security risk."** Obfuscation plus fail-closed disclosure: only measured, permitted material is ever exposed; secrets, credentials, and the moat never leave. Safety is a product guarantee, not a promise.
- **"My code is a proprietary asset."** Most of it is a maintenance liability, not an asset. And co-ownership means you keep it and keep using it — you are licensing measured value, not selling exclusivity.

- **“Who would buy plumbing?”** Labs, badly. Ordinary production code is exactly a coding-agent's textbook — and it is rights-clean, which scraped code is not. Boring to you is training gold to them.
- **“Reputation.”** Framed correctly, it is monetizing idle non-core assets — ordinary capital allocation — and it can be fully private and obfuscated. Companies already give this class of code away for free (see below); selling it is less exposure, not more.

2.5 The proof already exists — open source

The strongest evidence the premise is true: companies already release non-moat code **publicly, for free**, for strategic benefit — React, Kubernetes, and thousands of internal tools. If giving non-moat code away is safe and routine (the well-worn “commoditize your complement” strategy), then *selling* it — while exposing *less* (obfuscated) and keeping ownership (co-ownership) — is a strict improvement. Bitcode simply monetizes what firms already prove is safe to release.

2.6 Why Bitcode makes the thesis actionable, not just true

The idea has been un-actionable because there was no way to price non-moat code, expose it safely, sell it without losing it, or reach buyers — and because deciding what is moat is hard. Bitcode solves each: **measurement** prices it; **obfuscation + fail-closed** expose only the safe parts; **co-ownership** keeps it yours; the **demand side** aggregates labs and agents; and **need-fit measurement** does the moat/non-moat sorting for you — you set the obfuscation boundary, the market prices the rest. A theoretically-true idea becomes a one-click revenue line.

3. Why IP-recycling is the most defensible supply

Andreessen Horowitz's analysis of data moats is decisive: durable advantage comes **not from volume** — most “data network effects” are scale effects that erode, because the next data point costs more and adds less — but from **exclusive sourcing rivals cannot scrape or buy**, verified provenance, a measurement feedback loop, and switching costs. IP-recycling supply is the **only pool that is genuinely exclusive**: a company's private code exists nowhere else to be scraped or bought. Paired with Bitcode's measurement (provenance + quality) and co-ownership / \$BTD lock-in, it is the most defensible supply on the board — and the largest and highest-quality. Its sole weakness is unlock difficulty: the reflex to keep all code private. That reflex is exactly what this thesis dismantles.

4. Attraction • competition • defensibility • speed

Attracting supply. Measurement turns idle code into priced, listed inventory in clicks; obfuscation makes deposit safe; co-ownership + \$BTD let sellers keep ownership and earn (including on read-synthesized packs); zero-friction self-serve unlocks the long tail; automated acceptance takes all comers.

Competing. Incumbents (Scale, Surge, Mercor) *manufacture* supply with rented human labor — expensive, mercenary, multi-homed, and visibly wage-pressured (Mercor cut pay $\sim 1/3$, \$21→\$16; Scale lost Google, OpenAI, and xAI within hours of Meta's 49% stake). Each incremental expert costs *more*, not less. Bitcode taps *existing* code — cheaper, larger, compounding, and a different pool entirely. It never fights for contractors; it unlocks the world's idle code.

Defensibility. Exclusive sourcing (above all IP-recycling) + measured provenance + co-ownership lock-in + the Depository's data-network-effect (more deposits → better read-synthesis → more buyers → more sellers). The moat strengthens over time, with a compliance tailwind (EU AI Act) that scraped and manufactured supply cannot match.

Speed. Existing code needs no creation lag, and permissionless self-serve supply has historically scaled **1–3 orders of magnitude** faster and cheaper than curated recruiting (GitHub added 36M developers in a year; token networks bootstrapped millions of contributors in 6–18 months). Research-for-cash ignites fastest to seed the two-sided market; IP-recycling is the slower, far bigger land-grab; defunct-IP is the uncontested flank. The unlock that makes “accept everything” safe is an **automated measurement gate** — every open platform pays a spam/quality tax (Hugging Face's tens of thousands of unsafe models, Steam's 80% shovelware), and the winners solve it with an automated trust signal, not human review. Objectively measurable code — verifiable by execution and tests — is the ideal substrate for exactly that gate.

5. What would make the thesis bulletproof (research to run)

- **Codebase-composition study.** Quantify the non-moat share across real enterprise codebases — target: show the safely-sellable fraction is large, and name the categories (infra, tooling, integrations, dead code).
- **Design-partner pilots.** Two or three companies deposit obfuscated non-moat repos; publish realized revenue and the zero-exposure result as ROI case studies.

- **A moat-safety rubric.** A simple, credible checklist for what is safe to deposit, so the people who must approve it (CTOs, security, legal) have a defensible framework.
- **Harden the two technical risks.** Measurement-gaming resistance and rights-integrity at deposit, each with verifiable, on-chain proofs.

In one line

The incumbents rent labor that never compounds. Bitcode taps the world's idle code and prices it by machine — accepting all supply permissionlessly while an automated measurement gate carries the trust. Among the pools, IP-recycling is the heaviest prize: the largest, highest-quality, only-exclusive supply on Earth — and the “sell your non-moat code” thesis, once proven, converts every successful company's biggest idle asset into supply no competitor can reach.

Advanced Engineered Software, Inc. — July 2026 · Sources: internal Bitcode Market Analysis & Supply-Demand briefs; a16z “The Empty Promise of Data Moats”; NFX marketplace tactics; GitHub Octoverse; industry technical-debt estimates.